

TABLE OF CONTENTS

Sl. No.	Contents	Page No.
1.	Role of each student	1
2.	Data Set	1
3.	Description of Project	1
5.	Problem Statement	2
6.	Objective, Approach Adopted	2
7.	Program Source Code	3-22
8.	Result and Conclusion	23

Mobile Price Classification

ROLE OF EACH STUDENT :

1). **Narendra Rathod** - Analysis of Dataset, Splitting of Train, Model Training using SVM Linear Kernel Classifier, Model Training using Gradient Boosting Classifier, Bias and Variance Calculations, Confusion Matrix.

2). **Nayan Kumar** - Test dataset, Feature Scaling, Model Training using KNN Classifier, Model Training XGBClassifier, Model Training using Ada Boost Classifier , Theory and Accuracy.

DATA SET

This dataset contain information about many mobiles and variables about it.

DATA SOURCE LINK - [CLICK ON THE LINK](#)

DESCRIPTION OF PROJECT:

Price is the most effective attribute of marketing and business. The very first question of costumer is about the price of items. All the costumers are first worried and thinks “If he would be able to purchase something with given specifications or not”.

The price of a product is the most important attribute of marketing that product. One of those products where price matters a lot is a smartphone because it comes with a lot of features so that a company thinks a lot about how to price this mobile which can justify the features and also cover the marketing and manufacturing costs of the mobile. In this article, I will walk you through the task of mobile price classification with Machine Learning using Python.

PROBLEM STATEMENT:

Narendra has started his own mobile company. He wants to give tough fight to big companies like Apple, Samsung etc.

He does not know how to estimate price of mobiles his company creates. In this competitive mobile phone market you cannot simply assume things. To solve this problem, he collects sales data of mobile phones of various companies.

Narendra wants to find out some relation between features of a mobile phone(eg:- RAM,Internal Memory etc) and its selling price. But he is not so good at Machine Learning. So he needs your help to solve this problem.

In this problem you do not have to predict actual price but a price range indicating how high is the price ?

OBJECTIVE:

The objective of the data source is to classify activities into one of the three activities performed.

The three activities performed were as follows:

- Cleaning Data and Feature Selection
- Training and Evaluate Model
- Save Model for Future Use

APPROACH ADOPTED:

With the given data set to get better idea of mobile price, we adopted various classifications/clustering algorithms and try to find out which classifier gives the best accuracy/score according to given data.

PROGRAM SOURCE CODE:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.model_selection import GridSearchCV
import warnings
warnings.filterwarnings('ignore')
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

df=pd.read_csv("train.csv")
df.head(10)
```

m_dep \ battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0 0.6	842	0	2.2	0	1	0	7
1 0.7	1021	1	0.5	1	0	1	53
2 0.9	563	1	0.5	1	2	1	41
3 0.8	615	1	2.5	0	0	0	10
4 0.6	1821	1	1.2	0	13	1	44
5 0.7	1859	0	0.5	1	3	0	22
6 0.8	1821	0	1.7	0	4	1	10
7 0.8	1954	0	0.5	1	0	0	24
8 0.7	1445	1	0.5	0	0	0	53
9 0.1	509	1	0.6	1	2	1	9

mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
talk_time \							
0	188	2	...	20	756	2549	9
19							7

1	136	3	...	905	1988	2631	17	3
7								
2	145	5	...	1263	1716	2603	11	2
9								
3	131	6	...	1216	1786	2769	16	8
11								
4	141	2	...	1208	1212	1411	8	2
15								
5	164	1	...	1004	1654	1067	17	1
10								
6	139	8	...	381	1018	3220	13	8
18								
7	187	4	...	512	1149	700	16	3
5								
8	174	7	...	386	836	1099	17	1
20								
9	93	5	...	1137	1224	513	19	10
12								

	three_g	touch_screen	wifi	price_range
0	0	0	1	1
1	1	1	0	2
2	1	1	0	2
3	1	0	0	2
4	1	1	0	1
5	1	0	0	1
6	1	0	1	3
7	1	1	1	0
8	1	0	0	0
9	1	0	0	0

[10 rows x 21 columns]

df.shape

(2000, 21)

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power   2000 non-null   int64
1   blue            2000 non-null   int64
2   clock_speed     2000 non-null   float64
3   dual_sim        2000 non-null   int64
4   fc              2000 non-null   int64
5   four_g          2000 non-null   int64
6   int_memory      2000 non-null   int64
```

```

7    m_dep          2000 non-null float64
8    mobile_wt      2000 non-null int64
9    n_cores        2000 non-null int64
10   pc             2000 non-null int64
11   px_height      2000 non-null int64
12   px_width       2000 non-null int64
13   ram            2000 non-null int64
14   sc_h           2000 non-null int64
15   sc_w           2000 non-null int64
16   talk_time      2000 non-null int64
17   three_g        2000 non-null int64
18   touch_screen   2000 non-null int64
19   wifi           2000 non-null int64
20   price_range    2000 non-null int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB

```

```
df.isnull().sum()
```

```

battery_power    0
blue             0
clock_speed      0
dual_sim         0
fc              0
four_g          0
int_memory       0
m_dep           0
mobile_wt       0
n_cores         0
pc              0 px_height
0 px_width      0
ram             0
sc_h            0
sc_w            0
talk_time       0
three_g         0
touch_screen    0
wifi            0
price_range     0
dtype: int64

```

```
df.describe()
```

	battery_power	blue	clock_speed	dual_sim	fc
\count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500
std	439.418206	0.5001	0.816004	0.500035	4.341444

min	501.000000	0.0000	0.500000	0.000000	0.000000
25%	851.750000	0.0000	0.700000	0.000000	1.000000
50%	1226.000000	0.0000	1.500000	1.000000	3.000000
75%	1615.250000	1.0000	2.200000	1.000000	7.000000
max	1998.000000	1.0000	3.000000	1.000000	19.000000

	four_g	int_memory	m_dep	mobile_wt	n_cores
... \					
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	0.521500	32.046500	0.501750	140.249000	4.520500
std	0.499662	18.145715	0.288416	35.399655	2.287837
min	0.000000	2.000000	0.100000	80.000000	1.000000
25%	0.000000	16.000000	0.200000	109.000000	3.000000
50%	1.000000	32.000000	0.500000	141.000000	4.000000
75%	1.000000	48.000000	0.800000	170.000000	7.000000
max	1.000000	64.000000	1.000000	200.000000	8.000000
...					

	px_height	px_width	ram	sc_h	sc_w
\					
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	645.108000	1251.515500	2124.213000	12.306500	5.767000
std	443.780811	432.199447	1084.732044	4.213245	4.356398
min	0.000000	500.000000	256.000000	5.000000	0.000000
25%	282.750000	874.750000	1207.500000	9.000000	2.000000
50%	564.000000	1247.000000	2146.500000	12.000000	5.000000
75%	947.250000	1633.000000	3064.500000	16.000000	9.000000
max	1960.000000	1998.000000	3998.000000	19.000000	18.000000

	talk_time	three_g	touch_screen	wifi
price_range				
count	2000.000000	2000.000000	2000.000000	2000.000000
mean	11.011000	0.761500	0.503000	0.507000
std	5.463955	0.426273	0.500116	0.500076
min	2.000000	0.000000	0.000000	0.000000
25%	6.000000	1.000000	0.000000	0.000000
50%	11.000000	1.000000	1.000000	1.000000
75%	16.000000	1.000000	1.000000	1.000000
max	20.000000	1.000000	1.000000	1.000000

[8 rows x 21 columns]

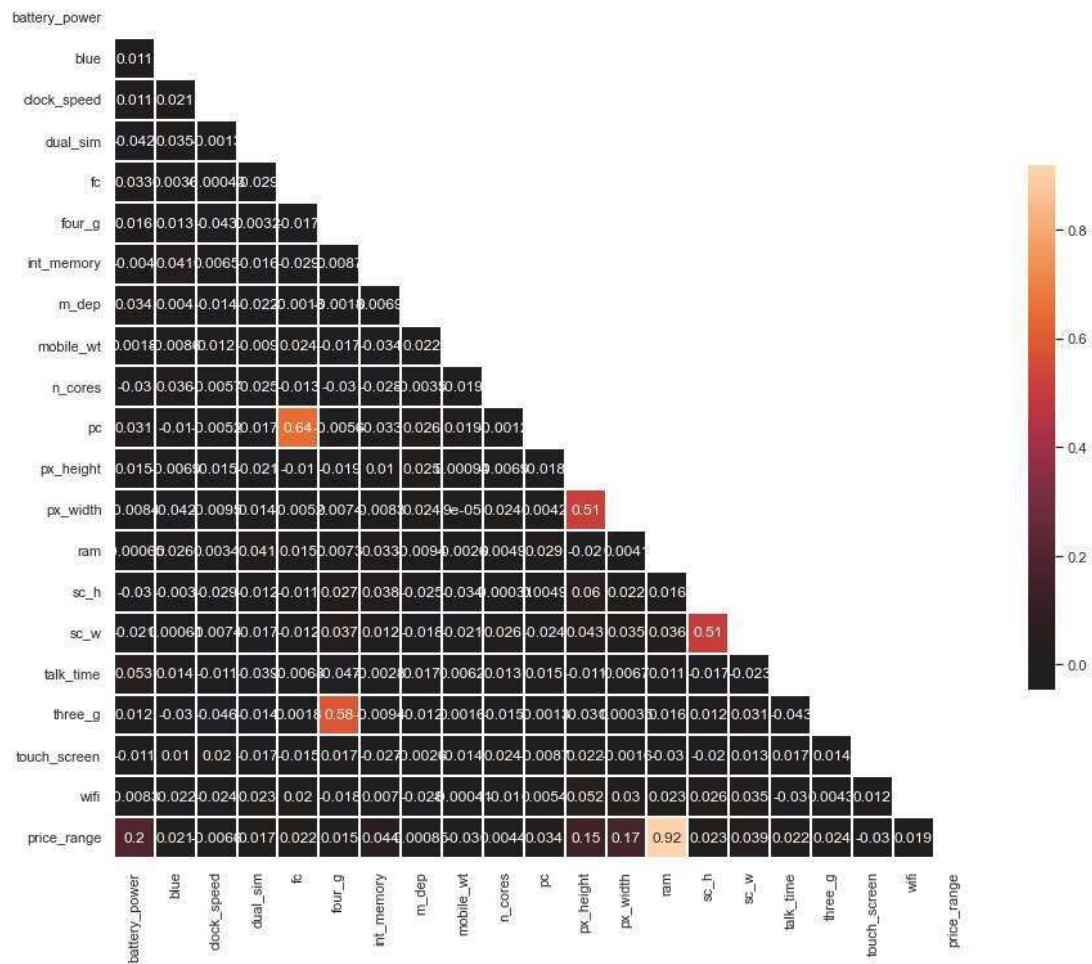
```
sns.set_theme(style="white")
```

```
corr = df.corr()
mask = np.triu(np.ones_like(corr, dtype=bool))
```


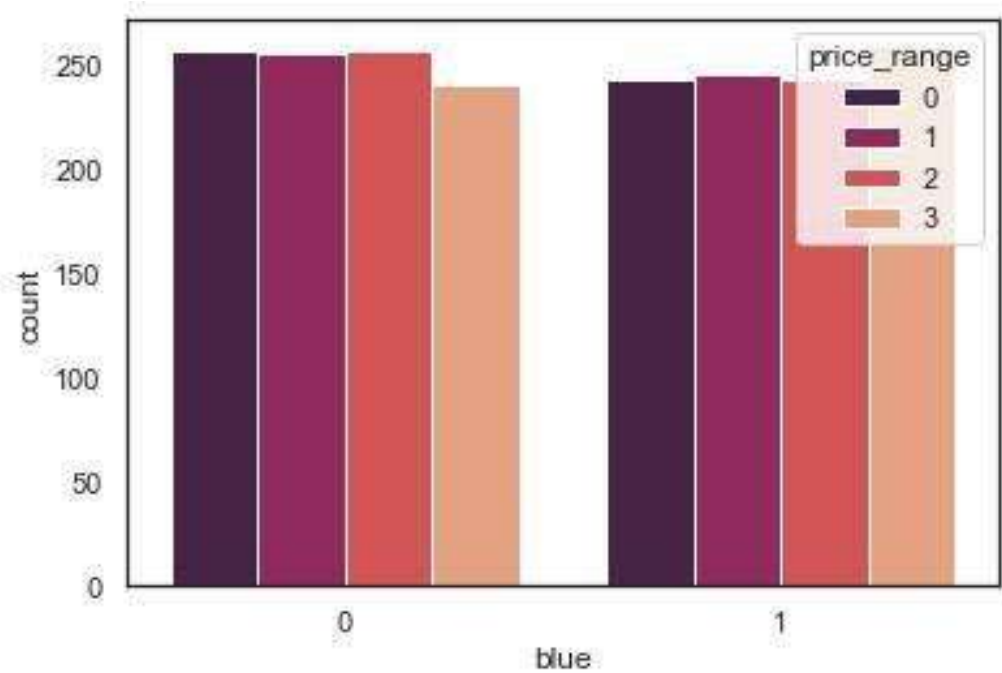
```
f, ax = plt.subplots(figsize=(15, 15))
cmap = sns.diverging_palette(230, 20)
```



```
sns.heatmap(corr, mask=mask, center=0, annot=True,
            square=True, linewidths=.3, cbar_kws={"shrink": 0.5})
```

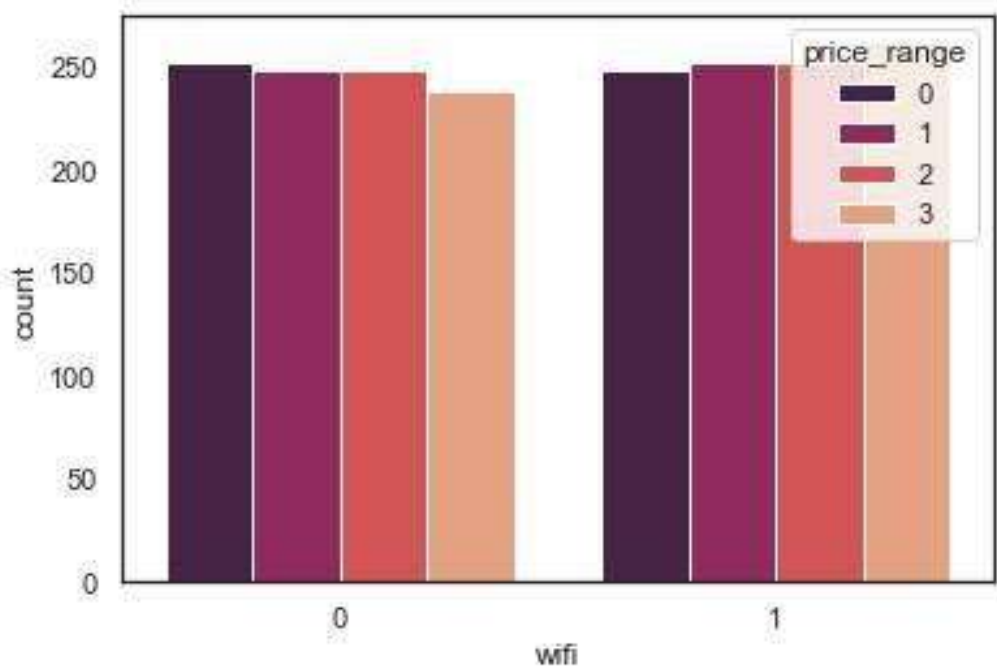
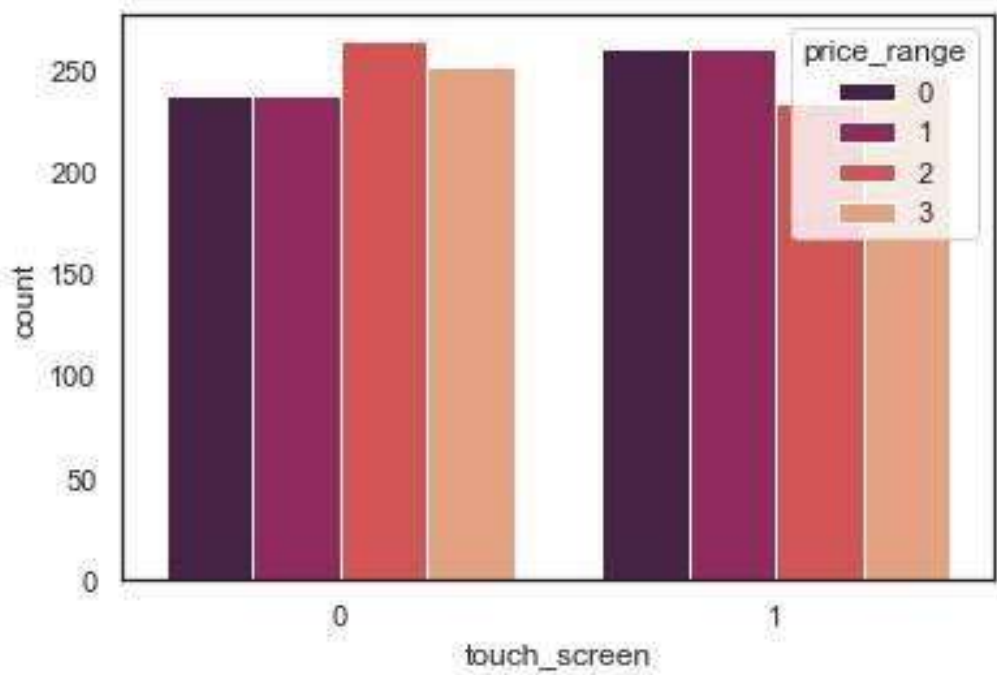
<AxesSubplot:>




```
binary_col =
['blue', 'dual_sim', 'four_g', 'three_g', 'touch_screen', 'wifi']
for i in binary_col:
    ax = sns.countplot(x=i,
hue='price_range', data=df, palette='rocket')
plt.show()
```









```
sns.displot(data=df, x="ram", hue="price_range", col="price_range")  
plt.show()
```




```
sns.pointplot(x='price_range',y='ram',data=df)
plt.show()
```



```
sns.boxplot(x='price_range',y='ram',data=df)
plt.show()
```




```
plt.figure(figsize=(12,6)) plt.subplot(1,2,1)
sns.barplot(x='price_range',y='px_height',data=df,palette='Blues')
plt.subplot(1,2,2)
sns.barplot(x='price_range',y='px_width',data=df,palette='Blues')
plt.show()
```




```
plt.figure(figsize=(12,6)) plt.subplot(1,2,1)
sns.barplot(x='price_range',y='fc',data=df,palette='Reds')
plt.subplot(1,2,2)
```


```
sns.barplot(x='price_range',y='pc',data=df,palette='Reds')
plt.show()
```




```
sns.pointplot(x='price_range',y='battery_power',data=df)
plt.show()
```



```
sns.set_style('darkgrid') sns.set_palette('Set2')
sns.barplot(x='price_range',y='battery_power',data=df)
plt.show()
```



```
sns.boxplot(x='price_range',y='battery_power',data=df)
<AxesSubplot:xlabel='price_range', ylabel='battery_power'>
```



```
f, ax = plt.subplots(figsize=(14, 14))
plt.subplot(1,2,1)
ax=sns.swarmplot(x="four_g", y="ram", hue="price_range",
                 palette="Dark2", data=df)
```




```
ax=sns.set(style="darkgrid")
```


```
plt.subplot(1,2,2)
```

```
ax=sns.swarmplot(x="three_g", y="ram", hue="price_range",  
                palette="Dark2", data=df)
```

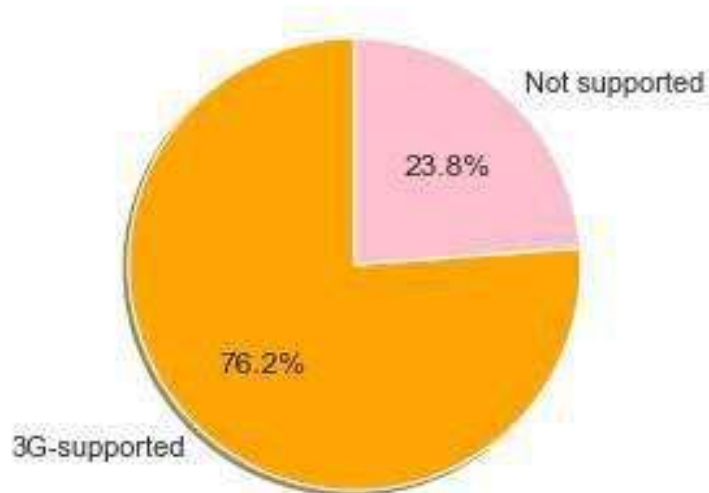
```
ax=sns.set(style="darkgrid")
```



```
fig1, ax1 = plt.subplots()
columns = ["4G-supported", 'Not supported']
ax1.pie(df.four_g.value_counts().values, labels=columns,
        autopct='%1.1f%%', shadow=True, startangle=90)
plt.show()
```



```
fig2, ax1 = plt.subplots()
columns = ["3G-supported", 'Not supported'] colors =
['orange', 'pink']
ax1.pie(df.three_g.value_counts().values, labels=columns,
autopct='%1.1f%%', shadow=True, startangle=90, colors=colors)
plt.show()
```



```
x=df.drop(['price_range'],axis=1)
y=df['price_range']

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=23)
```

```

parameters = {'n_neighbors': np.arange(1,20,1)}
knn=KNeighborsClassifier()
clf = GridSearchCV(knn, parameters)
clf.fit(x_train,y_train)

GridSearchCV(estimator=KNeighborsClassifier(),
              param_grid={'n_neighbors': array([ 1,  2,  3,  4,  5,  6,
 7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,
 18, 19])})

knn=KNeighborsClassifier(**clf.best_params_)
knn.fit(x_train,y_train)

KNeighborsClassifier(n_neighbors=15)

y_pred=knn.predict(x_test)
print("\nConfusion Matrix:\n%s"%confusion_matrix(y_pred,y_test))
print("\nTest Set Accuracy:"+str(accuracy_score(y_pred,y_test)*100))

```

```

Confusion Matrix:
[[148  6  0  0]
 [ 3 141  6  0]
 [ 0  7 134  5]
 [ 0  0  8 142]]

```

Test Set Accuracy:94.16666666666667

```

knn_acc=accuracy_score(y_pred,y_test)
knn_acc

```

0.9416666666666667

```

svc = SVC()
svc.fit(x_train, y_train)

```

SVC()

```

y_pred=svc.predict(x_test)
print("\nConfusion Matrix:\n%s"%confusion_matrix(y_pred,y_test))
print("\nTest Set Accuracy:"+str(accuracy_score(y_pred,y_test)*100))

```

```

Confusion Matrix:
[[148  5  0  0]
 [ 3 143  3  0]
 [ 0  6 138  0]
 [ 0  0  7 147]]

```

Test Set Accuracy:96.0

```
svc_acc=accuracy_score(y_pred,y_test)
svc_acc
```

0.96

```
parameters = {'booster':'gbtree',
              'learning_rate':[0.1,0.2,0.3],
              'max_depth':[3,4,5],
              'n_estimators':[5,8,10],
              'gamma':[3,4,5]}
```

```
xgb=XGBClassifier()
clf = GridSearchCV(xgb, parameters)
clf.fit(x_train,y_train,eval_metric='rmse')
```

```
GridSearchCV(estimator=XGBClassifier(base_score=None, booster=None,
                                     colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None,
                                     enable_categorical=False,
                                     gamma=None,
                                     gpu_id=None,
                                     importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=None,
                                     max_delta_step=None,
                                     max_depth=None,
                                     min_child_weight=None,
                                     missing=nan,
                                     monotone_constraints=None,
                                     n_estimators=100, n_jobs=None,
                                     num_parallel_tree=None,
                                     predictor=None,
                                     random_state=None,
                                     reg_alpha=None,
                                     reg_lambda=None,
                                     scale_pos_weight=None,
                                     subsample=None, tree_method=None,
                                     validate_parameters=None,
                                     verbosity=None),
```

```
        param_grid={'gamma': [3, 4, 5], 'learning_rate': [0.1,
0.2, 0.3],
                    'max_depth': [3, 4, 5], 'n_estimators': [5,
8, 10]})
```

```
xgb = XGBClassifier(**clf.best_params_)
xgb.fit(x_train,y_train,eval_metric='rmse')
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1,
              enable_categorical=False,
              gamma=3, gpu_id=-1, importance_type=None,
```

```

        interaction_constraints='', learning_rate=0.3,
max_delta_step=0,
        max_depth=5, min_child_weight=1, missing=nan,
        monotone_constraints='()', n_estimators=10, n_jobs=4,
        num_parallel_tree=1, objective='multi:softprob',
predictor='auto',
        random_state=0, reg_alpha=0, reg_lambda=1,
scale_pos_weight=None,
        subsample=1, tree_method='exact', validate_parameters=1,
        verbosity=None)

```

```

y_pred=xgb.predict(x_test)
print("\nConfusion Matrix:\n%s"%confusion_matrix(y_pred,y_test))
print("\nTest Set Accuracy:"+str(accuracy_score(y_pred,y_test)*100))

```

```

Confusion Matrix:
[[141  11   0   0]
 [ 10 126  11   0]
 [   0  17 123   9]
 [   0   0  14 138]]

```

Test Set Accuracy:88.0

```

xgb_acc=accuracy_score(y_pred,y_test)
xgb_acc

```

0.88

```

gbc = GradientBoostingClassifier()
gbc.fit(x_train, y_train)

```

```

GradientBoostingClassifier()

```

```

y_pred=gbc.predict(x_test)
print("\nConfusion Matrix:\n%s"%confusion_matrix(y_pred,y_test))
print("\nTest Set Accuracy:"+str(accuracy_score(y_pred,y_test)*100))

```

```

Confusion Matrix:
[[146  10   0   0]
 [   5 128   8   0]
 [   0  16 135   8]
 [   0   0   5 139]]

```

Test Set Accuracy:91.33333333333333

```

gbc_acc=accuracy_score(y_pred,y_test)
gbc_acc

```

0.9133333333333333

```

parameters = {'learning_rate':[0.01,0.1,1,10],
              'n_estimators':[50,80,100,120,150,180],
              'algorithm':['SAMME', 'SAMME.R']}
abc=AdaBoostClassifier()
clf = GridSearchCV(abc, parameters)
clf.fit(x_train,y_train)

GridSearchCV(estimator=AdaBoostClassifier(),
              param_grid={'algorithm': ['SAMME', 'SAMME.R'],
                          'learning_rate': [0.01, 0.1, 1, 10],
                          'n_estimators': [50, 80, 100, 120, 150,
180]}))

abc = AdaBoostClassifier(**clf.best_params_)
abc.fit(x_train, y_train)

AdaBoostClassifier(algorithm='SAMME', learning_rate=0.1,
n_estimators=150)

y_pred=abc.predict(x_test)
print("\nConfusion Matrix:\n%s"%confusion_matrix(y_pred,y_test))
print("\nTest Set Accuracy:"+str(accuracy_score(y_pred,y_test)*100))

```

```

Confusion Matrix:
[[126  22   0   0]
 [ 25 110  28   0]
 [  0  22  75  11]
 [  0   0  45 136]]

```

Test Set Accuracy:74.5

```

abc_acc=accuracy_score(y_pred,y_test)
abc_acc

0.745

```

```

models = pd.DataFrame({
    'Model': ['KNN','SVC','XGBClassifier','Gradient Boosting
Classifier',
             'Ada Boost Classifier'],
    'Score': [knn_acc, svc_acc, xgb_acc, gbc_acc, abc_acc]
})

```

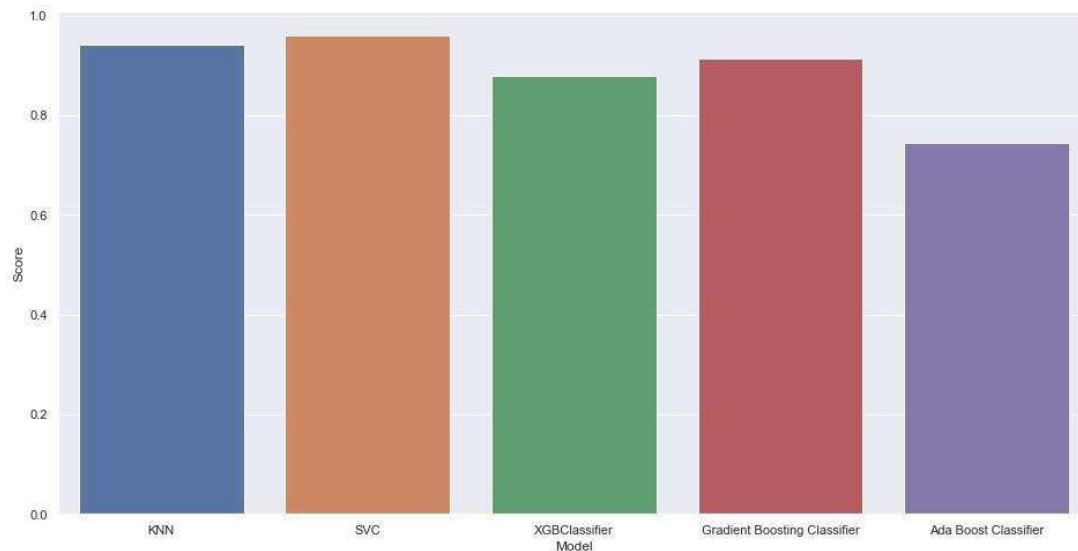
```

models.sort_values(by = 'Score', ascending = False)

```

	Model	Score
1	SVC	0.960000
0	KNN	0.941667
3	Gradient Boosting Classifier	0.913333
2	XGBClassifier	0.880000
4	Ada Boost Classifier	0.745000

```
plt.figure(figsize=(16,8))
sns.barplot(x='Model',y='Score',data=models)
<AxesSubplot:xlabel='Model', ylabel='Score'>
```



```
from mlxtend.evaluate import bias_variance_decomp

mse, bias, var = bias_variance_decomp(svc,x_train.values,
y_train.values, x_test.values, y_test.values, loss='mse',
num_rounds=100, random_seed=53)

print('Bias: %.3f' %bias)
Bias: 0.034

print('Variance: %.3f' %var)
Variance: 0.017

print('MSE: %.3f' %mse)
MSE: 0.051
```